

<div>on RequestVote request from peer</div> <pre> if currentTerm < m.term: stepDown(m.term) if (currentTerm == m.term and votedFor in [None, peer] and (m.lastLogTerm > logTerm(len(log)) or (m.lastLogTerm == logTerm(len(log)) and m.lastLogIndex >= len(log))): granted := True votedFor = peer electionAlarm = now() + rand(1.0, 2.0) * ELECTION_TIMEOUT else: granted := False reply {term: currentTerm, granted: granted} </pre> <div>on RequestVote response from peer</div> <pre> if currentTerm < m.term: stepDown(m.term) if (state == CANDIDATE and currentTerm == m.term): rpcDue[peer] = INFINITY voteGranted[peer] = m.granted </pre> <div>on AppendEntries request from peer</div> <pre> if currentTerm < m.term: stepDown(m.term) if currentTerm > m.term: reply {term: currentTerm, success: False} else: leader = peer state = FOLLOWER electionAlarm = now() + rand(1.0, 2.0) * ELECTION_TIMEOUT success := (m.prevIndex == 0 or (m.prevIndex <= len(log) and log[m.prevIndex].term == m.prevTerm)) if success: index := m.prevIndex for j := 1..len(m.entries): index += 1 if getTerm(index) != m.entries[j].term: log = log[1..(index-1)] + m.entries[j] commitIndex = min(m.commitIndex, index) else: index = 0 reply {term: currentTerm, success: success, matchIndex: index} </pre> <div>on AppendEntries response from peer</div> <pre> if currentTerm < m.term: stepDown(m.term) elif state == LEADER and currentTerm == m.term: if m.success: matchIndex[peer] = m.matchIndex nextIndex[peer] = m.matchIndex + 1 else: nextIndex[peer] = max(1, nextIndex[peer] - 1) </pre> <div>on StateMachine request from client</div> <pre> if state == LEADER: log.append({term: currentTerm, command: m.command}) </pre> <div>helper functions</div> <pre> def stepDown(newTerm): currentTerm = newTerm state = FOLLOWER votedFor = None if electionAlarm < now(): electionAlarm = now() + rand(1.0, 2.0) * ELECTION_TIMEOUT def logTerm(index): if index < 1 or index > len(log): return 0 else: return log[index].term </pre>	<div>start new election</div> <pre> on (state in [FOLLOWER, CANDIDATE] and electionAlarm < now()): electionAlarm = now() + rand(1.0, 2.0) * ELECTION_TIMEOUT currentTerm += 1 votedFor = serverID state = CANDIDATE foreach peer: # reset all state for peer </pre> <div>send RequestVote to peer</div> <pre> on (state == CANDIDATE and rpcDue[peer] < now()): rpcDue[peer] = now() + RPC_TIMEOUT send RequestVote to peer { term: currentTerm, lastLogTerm: logTerm(len(log)), lastLogIndex: len(log)} </pre> <div>become leader</div> <pre> on (state == CANDIDATE and sum(voteGranted) + 1 > NUM_SERVERS / 2: state = LEADER leader = localhost foreach peer: nextIndex[peer] = len(log) + 1 </pre> <div>send AppendEntries to peer</div> <pre> on (state == LEADER and (matchIndex[peer] < len(log) or rpcDue[peer] < now()): rpcDue[peer] = now() + ELECTION_TIMEOUT / 2 lastIndex := choose in (nextIndex[peer] - 1)..len(log) nextIndex[peer] = lastIndex send AppendEntries to peer { term: currentTerm, prevIndex: nextIndex[peer] - 1, prevTerm: getTerm(nextIndex[peer] - 1), entries: log[nextIndex[peer]..lastIndex], commitIndex: commitIndex} </pre> <div>advance commit index</div> <pre> n := sorted(matchIndex + [len(log)])[NUM_SERVERS / 2ish] on (state == LEADER and logTerm(n) == currentTerm): commitIndex = n </pre> <div>advance state machine</div> <pre> on lastApplied < commitIndex: lastApplied += 1 result := stateMachine.apply(log[lastApplied]) if (state == Leader and logTerm(lastApplied) == currentTerm): # send result to client </pre> <div>server state</div> <pre> # FOLLOWER, CANDIDATE, or LEADER state := FOLLOWER # latest term server has seen (increases monotonically) currentTerm := 1 # candidate that received vote in current term votedFor := None # log entries; each entry contains command for state # machine, and term when entry was received by leader log := [] # indexed from 1 # index of highest log entry known to be committed commitIndex := 0 # time after which to start new election electionAlarm := 0.0 # applies committed commands in log order stateMachine := new SM() # identity of last known leader leader := None # State per peer, valid only for the current term foreach peer: # time after which to send another RPC # (RequestVote or heartbeat) rpcDue[peer] := 0.0 # True if peer has granted this server its vote voteGranted[peer] := False # index of highest log entry known to be replicated # on peer matchIndex[peer] := 0 # index of next log entry to send to peer nextIndex[peer] := 1 </pre>
--	--

Figure 1: alternative cheatsheet